Rate and energy per token on 1, 2 and 4 boards

Each dot is one 400-token run; the open diamond is the same board without its engine (different software)

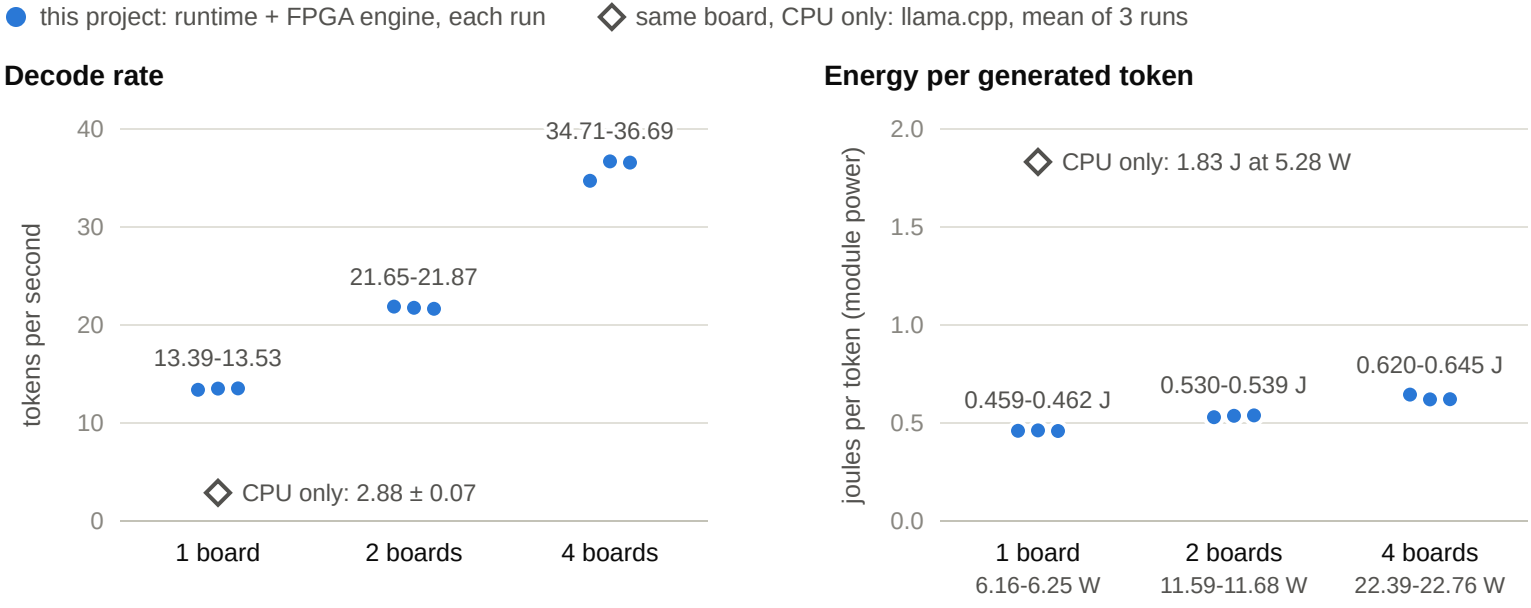

Llama-3.2-1B-Instruct, every weight matrix ggml Q4_0, 400 generated tokens. Boards: run-tp.sh, 24-token prompt (positions 24-423), greedy, 3 runs each; medians 13.51, 21.75, 36.57 tokens/s, 1.61x on two boards and 2.71x on four.
CPU only: llama-bench on board 1's four A53 cores, tg400 from an empty context (positions 0-399), 2.88 ± 0.07 (mean ± s.d.).
Power: the INA260 on each running K26 module, about 10 samples/s; carrier cards, fans, power supplies and the switch are not measured. 2026-11-04, buildlog 0013; source: evidence/0013-tp-runs.txt, -power-*, -cpu-llama-bench.txt.