Where one token's time goes

Time per generated token on board 1 (rank 0), three runs for each board count

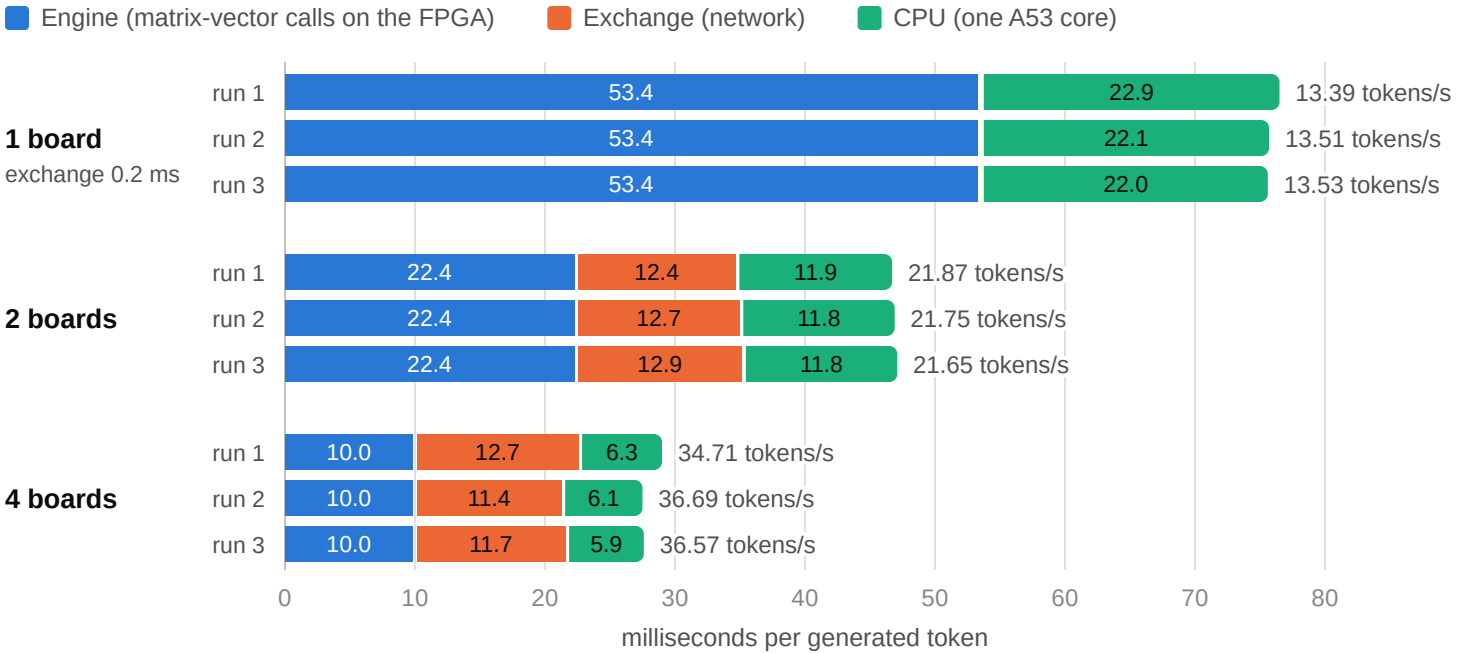

Llama-3.2-1B-Instruct, every weight matrix ggml Q4_0. run-tp.sh, 400 tokens generated after a 24-token prompt (positions 24-423), greedy; gemv2 bitstream at 250 MHz; rank 0's buckets: all 423 forwards, prompt included, over the 400 tokens (about 6% high). Engine time includes preparing each call's input on the CPU. 2026-11-04, buildlog 0013; source: evidence/0013-tp-runs.txt.